

Kenneth Willeford

CSCI 176

Final Project Report

Problem Description and Motivation

As brought up in the program documentation expectimax is an adversarial search algorithm that assumes near-optimal play(stochastic 'mistakes' are permitted.) it performs this by utilizing the idea of minimax (shown here with alpha-beta pruning heuristic - http://inst.eecs.berkeley.edu/~cs61b/fa14/ta-materials/apps/ab_tree_practice/) but instead of treating the values as absolutes it will average the benefit of actions. The importance of this algorithm can be demonstrated by visualizing the following...

The World is an infinite 1-Dimensional Grid

[| | | | | | | |]

The Protagonist Needs to move 3 spaces right to reach it's goal giving it a score of +9999

[| | | P | | G | | |]

However an antagonistic agent is between...

[| | | P | A | G | | |]

Each move the current score ticks down by 1, touching the antagonist ends the game with the current score.

Minimax would assume that A will play perfectly and P will decide the most optimal action is to suicide into A. Expectimax however doesn't assume that A is perfect and will leave things up to chance that A will move right twice giving P the chance to get a score of +9999(a big win.)

Unlike minimax expectimax doesn't have heuristics that are guaranteed to return the absolute best action (Alpha-Beta Pruning). Certain heuristics can provide pseudo-optimal play such as a "Better Than Last Time" heuristic where it is short-circuited when an action would put it in a better state than the action from a previous run of the algorithm did. But this ignores the chance for 'big wins' which is a strength of expectimax. For this reason parallelization is key to improving the performance of expectimax over a heuristic approach.

Baseline Solution

The general idea of expectimax is to determine the average utility of states producible by an action and to select the action with the highest average. A higher 'problem depth' essentially looks additional moves into the future. For example with a depth of 15 the algorithm is considering the next 15 moves that can be made. A more general implementation would consider if the next agents are protagonistic or antagonistic and control the flow between minimization and maximization respectively. For simplicity it is assumed there are two agents playing against each other. The implementation follows...

```
action* expectimax(state<action> s, int maxDepth){
    // Begin with score of negative infinity
    float bestScore = -9999;
    action* bestAction = NULL;
    for(int i = 0; i < NUM_ACTIONS; i++){
        // Get Each Action
        action* currAction = (s.Agents[s.CurrentAgent]).Actions[i];
        //if(!(currAction->Possible(s))){ continue; }
        // Get the utility of each action.
        float expmaxVal = expectiminValue(currAction,s, maxDepth);
        // Check to see if the current action is better than previous.
        if(expmaxVal > bestScore){
            bestScore = expmaxVal;
            bestAction = currAction;
        }
    }
    // return the best action to take.
    return bestAction;
}

float expectiminValue(action* currAction,state<action> s, int depth){
    // If end has been reached then evaluate the action and terminate.
    if(currAction->Evaluate(s) == 9999 || depth <= 0) return currAction->Evaluate(s);
    // Otherwise initialize the utility value as 0.
    float expVal = 0;
    // Get the next state.
    state<action> sS = currAction->Perform(s);
    for(int i = 0; i < NUM_ACTIONS; i++){
        // Get each of the successor actions..
        action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
```

```

        // if(!(nextAction->Possible(sS))) continue;
        // get the utility of the resulting states.
        expVal += expectimaxValue(nextAction,sS, depth-1);
    }
    // Average the utility among possible actions.
    return expVal / NUM_ACTIONS;
}
// Essentially repeats the main function call, the catch is that the score is what is cared about, not
the action.
float expectimaxValue(action* currAction,state<action> s, int depth){
    // If end has been reached then evaluate the action and terminate.
    if(currAction->Evaluate(s) == 9999 || depth <= 0) return currAction->Evaluate(s);
    // Otherwise initialize the best score as negative infinity.
    float bestScore = -9999;
    // Get the next state.
    state<action> sS = currAction->Perform(s);
    for(int i = 0; i < NUM_ACTIONS; i++){
        // Get each of the successor actions
        action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
        //if(!(nextAction->Possible(sS))) { continue; }
        // Get the utility of each of the successor states.
        float expminVal = expectiminValue(nextAction, sS, depth-1);
        // Whenever a better utility is found use it.
        if(expminVal > bestScore) bestScore = expminVal;
    }
    // return the found utility.
    return bestScore;
}

```

Parallel Solution and Implementation Methodology

The serial solution required some modification to allow for parallelization. For example values had to be utilized later in a way that reduction could not consider, for this reason the output variables were converted into output arrays. These loops also had only part of it's content able to be parallelized, but this wasn't a problem as most of the work being done in the loop could be parallelized. Breaking apart the loops and taking advantage of the output arrays allowed for a majority of the speedup. Another thing to consider is just how massive the problem size can become. Parallel For with only 1 thread would still incur some overhead which normally wouldn't be an issue, but in this case the problem size is $1 + \text{NUM_ACTIONS}^2 + \dots + \text{NUM_ACTIONS}^n$ with that small overhead applied on every step of the problem. For this reason the Parallel For is ignored when only 1 thread would be considered and instead at that phase it would be passed to the serial implementation. The parallel implementation follows...

```
int NeededThreads(int numThreads,int partitions){
    int output = (numThreads < NUM_ACTIONS) ? numThreads : NUM_ACTIONS;
    if(output <= 0){ output = 1; }
    return output;
}
action* expectimaxParallel(state<action> s, int maxDepth, int numThreads){
    action* bestAction = NULL;
    float expmaxVal[NUM_ACTIONS];

    /// First parallelizable section. This section provides the most benefit.
    /// A linear speedup should be possible until numThreads > NUM_ACTIONS
    int forThreads = NeededThreads(numThreads,NUM_ACTIONS);
    #pragma omp parallel for num_threads(forThreads)
    for(int i = 0; i < NUM_ACTIONS; i++){
        // Get Each Action
        action* currAction = (s.Agents[s.CurrentAgent]).Actions[i];
        // Get the utility of each action.
        expmaxVal[i] = expectiminValueParallel(currAction,s, maxDepth, numThreads -
forThreads);
    }
    // Begin with score of negative infinity
```

```

float bestScore = -9999;
/// Unparallelizable section is moved out. Computation of bestScore depends on previous
values.
for(int i = 0; i < NUM_ACTIONS; i++){
    // Check to see if the current action is better than previous.
    if(expmaxVal[i] > bestScore){
        bestScore = expmaxVal[i];
        bestAction = (s.Agents[s.CurrentAgent]).Actions[i];
    }
}
// return the best action to take.
return bestAction;
}
float expectiminValueParallel(action* currAction, state<action> s, int depth, int
remainingThreads){
    // If end has been reached then evaluate the action and terminate.
    if(currAction->Evaluate(s) == 9999 || depth <= 0) return currAction->Evaluate(s);
    // Otherwise initialize the utility value as 0.
    float expVal = 0;
    // Get the next state.
    state<action> sS = currAction->Perform(s);
    /// Second parallelizable section. Diminishing returns should be observed based on tree
depth.
    /// The explicit split is because of how many instances may be created, the overhead of
parallel_for on 1 thread was noticable.
    int forThreads = NeededThreads(remainingThreads, NUM_ACTIONS);
    if(forThreads > 1){
        #pragma omp parallel for num_threads(forThreads) reduction(+ : expVal)
        for(int i = 0; i < NUM_ACTIONS; i++){
            // Get each of the successor actions..
            action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
            // get the utility of the resulting states.
            expVal += expectimaxValueParallel(nextAction, sS, depth-1,
remainingThreads - forThreads);
        }
    } else {
        for(int i = 0; i < NUM_ACTIONS; i++){
            // Get each of the successor actions..
            action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
            // get the utility of the resulting states.
            expVal += expectimaxValueParallel(nextAction, sS, depth-1);
        }
    }
    // Average the utility among possible actions.
    return expVal / NUM_ACTIONS;
}

```

```

}
// Essentially repeats the main function call, the catch is that the score is what is cared about, not
the action.
float expectimaxValueParallel(action* currAction, state<action> s, int depth, int
remainingThreads){
    float expminVal[NUM_ACTIONS];
    // If end has been reached then evaluate the action and terminate.
    if(currAction->Evaluate(s) == 9999 || depth <= 0) return currAction->Evaluate(s);
    // Otherwise initialize the best score as negative infinity.
    float bestScore = -9999;
    // Get the next state.
    state<action> sS = currAction->Perform(s);
    /// Third parallelizable section. Diminishing returns should be observed based on tree
depth.
    /// The explicit split is because of how many instances may be created, the overhead of
parallel_for on 1 thread was noticable.
    int forThreads = NeededThreads(remainingThreads, NUM_ACTIONS);
    if(forThreads > 1){
        #pragma omp parallel for num_threads(forThreads)
        for(int i = 0; i < NUM_ACTIONS; i++){
            // Get each of the successor actions
            action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
            //if(!(nextAction->Possible(sS))){ continue; }
            // Get the utility of each of the successor states.
            expminVal[i] = expectiminValueParallel(nextAction, sS, depth-1,
remainingThreads);
        }
    } else {
        for(int i = 0; i < NUM_ACTIONS; i++){
            // Get each of the successor actions
            action* nextAction = (sS.Agents[sS.CurrentAgent]).Actions[i];
            //if(!(nextAction->Possible(sS))){ continue; }
            // Get the utility of each of the successor states.
            expminVal[i] = expectiminValue(nextAction, sS, depth-1);
        }
    }

    /// Second section which couldn't be parallelized. It is moved out.
    // Whenever a better utility is found use it.
    for(int i = 0; i < NUM_ACTIONS; i++)
        if(expminVal[i] > bestScore) bestScore = expminVal[i];
    // return the found utility.
    return bestScore;
}

```

Performance Comparison

The following runtime outputs are with num_threads = 1,2,4 respectively.(my system is an i5 so I don't get anymore benefit passed that.)

Compile With: g++ -fopenmp -o xxx final-project.cpp

Run With: ./xxx <num threads> <problem depth>

./xxx 1 12

Serial Execution Time of Expectimax with depth 12 is 49.91s.

Parallel Execution Time of Expectimax with depth 12 and 1threads is 51.046s.

./xxx 2 12

Serial Execution Time of Expectimax with depth 12 is 51.113s.

Parallel Execution Time of Expectimax with depth 12 and 2threads is 26.624s.

./xxx 4 12

Serial Execution Time of Expectimax with depth 12 is 53.786s.

Parallel Execution Time of Expectimax with depth 12 and 4threads is 14.397s.

Performance Analysis

The time complexity of expectimax is b^m , or the branching factor(NUM_ACTIONS in this case) to the depth that the problem is being taken. The speedup also depends largely on the branching factor. In my example the branching factor is 4. By parallelizing the top level with 4 threads we can the time taken is $T/4$. To receive a speedup on the next level we would need 4 threads per decrease in time taken. By parallelizing the next level using a total of 16 threads we can achieve $T/16$ (not counting overhead.) So it should be strongly scalable providing a linear speedup assuming that $p = b^n$ where n is any integer.