

Kenneth Willeford

CSCI 176

HW2

1. $T_{\text{serial}} = n$, $T_{\text{parallel}} = n/p + \log_2(p)$

If we increase p by factor of k

$$T_{\text{serial}} = n, T_{\text{parallel}} = n/(p*k) + \log_2(p*k)$$

If we are looking to maintain constant efficiency then $T_{\text{serial}} = T_{\text{parallel}}$ or

$$n/(p*k) + \log_2(p*k) = n$$

$$n - n/(p*k) = \log_2(p*k)$$

$$n*(p*k) - n = \log_2(p*k)*(p*k)$$

It is difficult to isolate n although wolfram finds a solution

[wolframalpha.com/input/?i=solve+for+n+in+n+-+n%2F\(p*k\)+%3D+log2\(p*k\)](http://wolframalpha.com/input/?i=solve+for+n+in+n+-+n%2F(p*k)+%3D+log2(p*k))

For my remaining purposes I will use it in this form.

“How much should we increase n by if we double the number of processes from 8 to 16?”

$$n_{\text{Increase}} = n_{\text{After}} - n_{\text{Before}}$$

$$n_{\text{After}}: 16n - n = 16*\log_2(16) ; 15n = 16*4 ; n = 64/15 = 4.27$$

$$n_{\text{Before}}: 8n - n = 8*\log_2(8) ; 7n = 8*3 ; n = 24/7 = 3.43$$

$$n_{\text{Increase}} = 4.27 - 3.43 = 0.84$$

The system requires an increase of input size in order for performance to not change. So it is scalable.

2. “Is a program that obtains linear speedup strongly scalable? Explain your answer.”

I would say “yes.” A linear speedup would imply an execution time of n/p which would mean that there are no diminishing returns on further parallelization. If this isn’t strongly scalable I don’t know what is.

3. Textbook 4.8

- A. This is the classic deadlock example. Deadlock occurs.
- B. With Busy-Waiting the issue would still occur. (Both are waiting on each other to finish up.)
- C. With a Semaphore the issue can be prevented if the semaphore allows up to two accesses to its synchronization variable. However it would be unsafe if the deadlock example scaled up.

4. Textbook 4.16

Left Operand: 8000x8000 Matrix

Right Operand: $X \times 8000$ Matrix where X is some arbitrary integer

Output Matrix: $X \times 8000$

4 threads: thread0 is on p0, thread2 is on p1

Cache Line = 64bytes (8 doubles)

Is it possible for false sharing to occur along vector y ? Why?

Yes, and it depends on the width of the output Matrix which is tied to X .

In order for there to be no false sharing between thread0 and thread2 there must be no way for their cache lines to overlap. There is 1 row separating the two. So we must find an overflow condition. If $X * 2 * 8\text{bytes} < 64\text{bytes}$ then false sharing will occur.

We can see this at the boundary case

□□□□ Thread0

□□□□ Thread1

□□□□ Thread2

If $64\text{bytes} < 64\text{bytes}$, the condition is false (and we can see that Thread0 just barely keeps out of thread 2.)

□□□□ Thread0

□□□□ Thread1

□□□□ Thread2

If $48\text{bytes} < 64\text{ bytes}$: The condition holds true, and sure enough if we check by hand Thread0's cache line is within Thread2's contents.

What if thread 0 and thread 3 are assigned to different processors, Is False Sharing Possible?

Yes, the condition is just adjusted slightly.

If $X * 3 * 8 \text{bytes} < 64 \text{bytes}$ then false sharing will occur.

Let's confirm with our boundary cases.

[][] Thread0

[][] Thread1

[][] Thread2

[][] Thread3

$2 * 3 * 8 \text{bytes} < 64 \text{bytes}$; $48 \text{bytes} < 64 \text{bytes}$, false sharing will occur. Checking by hand confirms.

[][][] Thread0

[][][] Thread1

[][][] Thread2

[][][] Thread3

$3 * 3 * 8 \text{bytes} < 64 \text{bytes}$; $72 \text{bytes} < 64 \text{bytes}$, false sharing will not occur. Checking by hand confirms.

5. Submitted Alongside

```

1  //////////////////////////////////////////////////
2  // Purpose: Estimate definite integral (or area under curve) using trapezoidal rule.
3  // - this version uses void type thread function, which uses global_result as ref
   para.
4  //
5  // Input:   a, b, n //assume n is evenly divisible by num_threads
6  // Output:  estimate of integral from a to b of f(x) using n trapezoids.
7  //
8  // Compile/run: $> g++ -fopenmp -o xxx hw2progl.cpp
9  //              $> xxx 4 //any number of threads
10 //
11 // Test Input      0 100000000 1000000
12 // Expected Output  3.33333e+023
13 //////////////////////////////////////////////////
14
15 #include <cstdlib>
16 #include <iostream>
17 #include <omp.h>
18 using namespace std;
19
20 double f(double); //f(x)=x^2
21 double Trap(double, double, int);
22
23 int main(int argc, char* argv[])
24 {
25     double global_result = 0.0;
26     double a, b; //left and right end points
27     int n; //total number of trapezoids
28
29     int thread_count = atoi(argv[1]); //command line arg
30     cout<<"Enter a, b, and n -- each separated by a space:\n";
31     cin>>a>>b>>n;
32     if(n % thread_count != 0)
33     { cerr<<"n should be evenly divisible by "<<thread_count<<endl;
34       exit(0);
35     }
36     #pragma omp parallel num_threads(thread_count)
37     global_result += Trap(a, b, n); //global_result is ref para
38
39     cout<<"With n="<<n<<" trapezoids, our estimate of integral from "
40         <<a<<" to "<<b<<" is "<<global_result<<endl;
41
42     return 0;
43 } /* main */
44
45 ////////////////////////////////////////////////// function f(x) = x * x ;
46 double f(double x)
47 {
48     return x*x;
49 }//f
50
51 ////////////////////////////////////////////////// thread function for Trap
52 // Purpose: Use trapezoidal rule to estimate definite integral
53 // Input args:
54 //     a, b : left, right endpoints
55 //     b: right endpoints
56 //     n: number of trapezoids
57 // Output arg (ref para):
58 //     globla_result: estimate of integral from a to b of f(x)
59 //////////////////////////////////////////////////
60 double Trap(double a, double b, int n)
61 {
62     double h, x, my_result;
63     double local_a, local_b;

```

```

64     int i, local_n;
65     int my_rank = omp_get_thread_num();
66     int thread_count = omp_get_num_threads(); //returns the count passed from pragma
parallel
67
68     h = (b-a)/n;
69     local_n = n/thread_count;
70     local_a = a + my_rank*local_n*h;
71     local_b = local_a + local_n*h;
72     my_result = (f(local_a) + f(local_b))/2.0;
73     for (i = 1; i <= local_n-1; i++)
74     { x = local_a + i*h;
75       my_result += f(x);
76     }
77     my_result = my_result*h;
78     #pragma omp critical
79     cout<<"Thread_"<<my_rank<<"", local_result = "<<my_result<<endl;
80
81     return my_result;
82 }//
83

```

```
1 Enter a, b, and n -- each separated by a space:
2 Thread_0, local_result = 3.33333e+023
3 With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
4
```

```
1 Enter a, b, and n -- each separated by a space:
2 Thread_1, local_result = 2.91667e+023
3 Thread_0, local_result = 4.16667e+022
4 With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
5
```

```
1 Enter a, b, and n -- each separated by a space:
2 Thread_1, local_result = 3.64583e+022
3 Thread_2, local_result = 9.89583e+022
4 Thread_3, local_result = 1.92708e+023
5 Thread_0, local_result = 5.20833e+021
6 With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
7
```



```

1  //////////////////////////////////////////////////
2  // Purpose: Estimate definite integral (or area under curve) using trapezoidal rule.
3  // - this version uses void type thread function, which uses global_result as ref
   para.
4  //
5  // Input:   a, b, n //assume n is evenly divisible by num_threads
6  // Output:  estimate of integral from a to b of f(x) using n trapezoids.
7  //
8  // Compile/run: $> g++ -fopenmp -o xxx hw2progl.cpp
9  //              $> xxx 4 //any number of threads
10 //
11 // Test Input      0 100000000 1000000
12 // Expected Output  3.33333e+023
13 //////////////////////////////////////////////////
14
15 #include <cstdlib>
16 #include <iostream>
17 #include <omp.h>
18 using namespace std;
19
20 double f(double); //f(x)=x^2
21 double Trap(double, double, int);
22
23 int main(int argc, char* argv[])
24 {
25     double global_result = 0.0;
26     double a, b; //left and right end points
27     int n; //total number of trapezoids
28
29     int thread_count = atoi(argv[1]); //command line arg
30     cout<<"Enter a, b, and n -- each separated by a space:\n";
31     cin>>a>>b>>n;
32     if(n % thread_count != 0)
33     { cerr<<"n should be evenly divisible by "<<thread_count<<endl;
34       exit(0);
35     }
36     #pragma omp parallel num_threads(thread_count) \
37         reduction(+: global_result)
38         global_result += Trap(a, b, n); //global_result is ref para
39     cout<<"With n="<<n<<" trapezoids, our estimate of integral from "
40         <<a<<" to "<<b<<" is "<<global_result<<endl;
41
42     return 0;
43 } /* main */
44
45 ////////////////////////////////////////////////// function f(x) = x * x ;
46 double f(double x)
47 {
48     return x*x;
49 }//f
50
51 ////////////////////////////////////////////////// thread function for Trap
52 // Purpose: Use trapezoidal rule to estimate definite integral
53 // Input args:
54 //     a, b : left, right endpoints
55 //     b: right endpoints
56 //     n: number of trapezoids
57 // Output arg (ref para):
58 //     globla_result: estimate of integral from a to b of f(x)
59 //////////////////////////////////////////////////
60 double Trap(double a, double b, int n)
61 {
62     double h, x, my_result;
63     double local_a, local_b;

```

```

64     int i, local_n;
65     int my_rank = omp_get_thread_num();
66     int thread_count = omp_get_num_threads(); //returns the count passed from pragma
parallel
67
68     h = (b-a)/n;
69     local_n = n/thread_count;
70     local_a = a + my_rank*local_n*h;
71     local_b = local_a + local_n*h;
72     my_result = (f(local_a) + f(local_b))/2.0;
73     for (i = 1; i <= local_n-1; i++)
74     { x = local_a + i*h;
75       my_result += f(x);
76     }
77     my_result = my_result*h;
78     #pragma omp critical
79     cout<<"Thread_"<<my_rank<<"", local_result = "<<my_result<<endl;
80
81     return my_result;
82 }//
83

```

```
1 Enter a, b, and n -- each separated by a space:
2 Thread_0, local_result = 3.33333e+023
3 With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
4
```

```
1 Enter a, b, and n -- each separated by a space:
2 Thread_0, local_result = 4.16667e+022
3 Thread_1, local_result = 2.91667e+023
4 With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
5
```

```
1  Enter a, b, and n -- each separated by a space:
2  Thread_0, local_result = 5.20833e+021
3  Thread_1, local_result = 3.64583e+022
4  Thread_2, local_result = 9.89583e+022
5  Thread_3, local_result = 1.92708e+023
6  With n=1000000 trapezoids, our estimate of integral from 0 to 1e+008 is 3.33333e+023
7
```